

Name of the Student	D.Chinna Moulali
Reg. No	192425278
GitHub Link	https://github.com/kiduuuu/MLA0405-Deep-learning

SIMATS ENGINEERING

CO4 - ASSESSMENT TOOL 2

Industry Problem Solving Task

COURSE NAME: Deep Learning
SLOT :

COURSE CODE: MLA0405

COURSE OUTCOME COVERED

CO 4: Students will be able to understand and apply probabilistic graphical models including Bayesian Networks, Markov Random Fields, Hidden Markov Models, and Conditional Random Fields. They will learn how to represent complex probabilistic relationships among variables and perform inference and learning in graphical models. Students will be able to use these models for reasoning under uncertainty and solving real-world decision-making problems. BL-6

Course Outcome Weightage: 10%
Assessment Tool Weightage: 40%

Duration: 90 minutes
Mark : 25

Code of Conduct:

I D.Chinna Moulali (192372384) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.
I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the student: D. Chinna Moulali

Name of the student: D.Chinna Moulali

Criteria	Understanding of Industry Problem (2)	Application of Engineering Concepts (2.5)	Solution Design (2.5)	Use of Tools (2)	Analysis (1)	Total (10)
Weightage	20 %	25%	25%	20 %	10 %	100%
Q1						
Q2						
Q3						
Q4						
Q5						
Total						

QUESTIONS: 05

Total Marks:25

Q1. Transfer Learning

A healthcare startup aims to develop an AI-based system for detecting skin cancer from dermoscopic images. Since only a limited number of labeled images are available, design a transfer learning strategy using a pre-trained CNN model. Justify your choice of model, identify which layers should be frozen or fine-tuned, and explain how your design improves classification performance.

Solution:

TRANSFER LEARNING FOR SKIN CANCER DETECTION

1. Industry Problem

A healthcare startup wants to develop an AI system that detects skin cancer from dermoscopic images. The major challenge is the limited availability of labeled medical images. Training a deep CNN from the beginning would require a very large dataset and significant computational resources.

Therefore, a transfer learning approach is suitable. A CNN pre-trained on a large image dataset can be adapted to dermoscopic images.

2. Proposed Model


I propose using EfficientNet-B0 as the pre-trained CNN backbone.


EfficientNet is suitable because:

- It provides strong image-classification performance.
- It achieves good accuracy with relatively fewer parameters.
- It is computationally efficient.
- It can extract useful low-level and high-level visual features.
- It is suitable for deployment in a healthcare startup where computational resources may be limited.

A suitable dataset is the ISIC (International Skin Imaging Collaboration) dataset, which contains dermoscopic skin-lesion images.

3. Architecture





4. Frozen and Fine-Tuned Layers

Initially, most of the EfficientNet layers should be frozen.

The early convolutional layers learn generic features such as:

- Edges
- Corners
- Textures
- Shapes

These features are useful for both natural images and medical images.

The final classification layer is removed and replaced with a new classification head designed for skin-cancer classes.

After initial training, the deeper convolutional layers can be unfrozen gradually and fine-tuned using a small learning rate.

Early layers → Frozen

Middle layers → Mostly frozen

Deep layers → Fine-tuned

Classification head → Newly trained

5. Training Strategy

The following preprocessing can be applied:

- Image resizing
- Pixel normalization
- Random rotation
- Horizontal/vertical flipping
- Zooming
- Random cropping
- Brightness variation

Data augmentation is important because medical datasets are usually limited.

Because malignant cases may be less frequent than benign cases, class-weighted loss or focal loss can be used to reduce class imbalance.

6. Performance Evaluation

The system should not be evaluated only using accuracy.

Important healthcare metrics include:

- Accuracy
- Precision
- Recall/Sensitivity
- Specificity
- F1-score
- ROC-AUC
- Confusion matrix

For cancer screening, sensitivity/recall is particularly important, because missing a malignant lesion can have serious consequences.

7. How Transfer Learning Improves Performance

Transfer learning provides the network with useful visual representations before training on the relatively small medical dataset.

Compared with training from scratch, it can:

- Reduce training time
- Reduce data requirements
- Improve generalization
- Reduce overfitting
- Improve classification performance
- Reduce computational requirements

8. Industry-Level Solution

The final system can be deployed as a clinical decision-support tool. The model should assist dermatologists rather than independently making a final medical diagnosis. A production system should also include external validation, calibration, bias testing across different skin types and imaging conditions, and monitoring after deployment.

Q2. DenseNet and PixelNet

A smart city project requires an automated road-scene analysis system that accurately identifies roads, pedestrians, vehicles, and traffic signs at the pixel level. Design a deep learning architecture using DenseNet and PixelNet to solve this problem. Explain how the proposed architecture improves feature reuse, segmentation accuracy, and computational efficiency.

Solution:

DENSENET AND PIXELNET FOR SMART-CITY ROAD SCENE ANALYSIS

1. Industry Problem

A smart-city system needs to understand road scenes at the **pixel level**.

The system should identify:


- Roads
- Pedestrians
- Vehicles
- Traffic signs
- Other relevant objects


Pixel-level classification is a semantic segmentation problem.

A normal image classifier only predicts a class for the complete image, whereas this application requires a class for every pixel.

2. Proposed Architecture

I propose combining a DenseNet feature extractor with a PixelNet-style pixel-wise prediction module. DenseNet is used to extract rich visual features, while PixelNet performs pixel-level classification.






3. DenseNet Feature Extraction

DenseNet uses dense connections.

Instead of each layer receiving information only from the previous layer, a layer receives feature maps from earlier layers.

Conceptually:




This encourages feature reuse.

For road-scene analysis, features such as edges, textures, vehicle shapes and road boundaries can be reused by later layers.

4. PixelNet Component

PixelNet-style processing converts the extracted feature representation into predictions for individual pixels.

For each pixel:



The output can contain a segmentation map where each pixel is assigned a class.

5. Feature Reuse

DenseNet improves feature reuse through dense connectivity.

Advantages include:

- Better gradient flow
- Reduced redundant feature learning
- Efficient feature propagation
- Reuse of low-level and high-level features
- Potentially fewer parameters than similarly deep conventional networks

6. Improving Segmentation Accuracy

To improve segmentation quality, the system can use:

- Multi-scale feature fusion
- Skip connections
- Bilinear interpolation or learned upsampling
- Dice loss
- Cross-entropy loss
- Data augmentation

- Boundary-aware loss

A combined segmentation loss can be written as:

Total Loss = Cross-Entropy Loss + λ (Dice Loss)

where λ controls the contribution of the Dice loss.

7. Computational Efficiency

DenseNet reduces redundant feature extraction through feature reuse. Pixel-wise classification can then operate on learned feature representations instead of independently processing every complete image region through a large network.

For deployment, the model can further be optimized using:

- Model pruning
- Quantization
- TensorRT
- GPU acceleration
- Efficient image resolution
- Batch processing

8. Dataset and Tools

Suitable datasets include:

- Cityscapes
- CamVid
- Mapillary Vistas

Tools:

- Python
- PyTorch or TensorFlow
- OpenCV
- CUDA
- GPU

9. Evaluation

The system can be evaluated using:

- Pixel accuracy
- Mean Intersection over Union (mIoU)
- Dice coefficient
- Precision
- Recall
- FPS
- Inference latency

mIoU is particularly useful for semantic segmentation.

10. Industrial Impact

The proposed system can support:

- Autonomous driving
- Smart traffic monitoring
- Pedestrian safety
- Traffic-sign recognition
- Road-condition monitoring
- Intelligent transportation systems

Therefore, combining DenseNet feature reuse with PixelNet-style pixel classification provides a practical approach for accurate road-scene segmentation while maintaining computational efficiency.

Q3. Bidirectional RNN and Encoder–Decoder

A multinational customer service company wants to build a multilingual chatbot capable of translating customer queries while preserving contextual meaning. Design an Encoder–Decoder sequence-to-sequence model using Bidirectional RNNs. Illustrate the architecture and justify how bidirectional processing improves translation accuracy.

Solution:

BIDIRECTIONAL RNN AND ENCODER-DECODER FOR MULTILINGUAL TRANSLATION


1. Industry Problem

A multinational customer-service company wants a chatbot that can translate customer queries between multiple languages while preserving contextual meaning.

For example:

Customer:


"I need to change my flight tomorrow."



The main challenge is that the meaning of a word can depend on surrounding words.

2. Proposed Architecture

I propose a Bidirectional RNN Encoder + Autoregressive RNN Decoder architecture.





Translated Sentence

An LSTM or GRU can be used instead of a simple RNN to improve long-sequence handling.

3. Bidirectional Processing

A conventional RNN reads:

I → need → to → change → my → flight

A bidirectional RNN processes the sequence in both directions:

Forward:

I → need → to → change → my → flight

Backward:

flight → my → change → to → need → I

The representation of each word therefore contains information from both its left and right context.

4. Why This Improves Translation

Consider:

"I booked a bank flight."

The surrounding context helps determine the intended meaning of words.


Bidirectional processing allows the encoder to capture:

- Previous context
- Future context
- Word relationships
- Sentence structure
- Long-range dependencies

This produces a richer context representation for the decoder.

5. Encoder


The input sentence is converted into tokens.



The encoder transforms the source-language sentence into meaningful hidden representations.

6. Decoder

The decoder generates the translated sentence one word at a time.



At every step, the decoder uses its hidden state and encoder information to predict the next word.

7. Attention Mechanism

For longer sentences, an attention mechanism can be added.


Attention allows the decoder to focus on the most relevant encoder states while generating each output word.

This helps prevent information loss caused by compressing an entire sentence into one vector.

8. Multilingual Extension

A language-ID token can be added to indicate the desired target language.

For example:



The same model can therefore support multiple language pairs.

9. Tools and Dataset

Possible datasets:

- Multi30K
- OPUS
- WMT datasets

Tools:

- Python
- PyTorch / TensorFlow
- Hugging Face libraries
- GPU/CUDA

10. Evaluation

Translation performance can be evaluated using:

- BLEU
- ROUGE
- METEOR
- Translation accuracy
- Human evaluation

The system should also be evaluated for contextual correctness rather than relying only on word-level accuracy.

11. Industry Benefit

The chatbot can provide:

- Multilingual customer support
- Real-time translation
- Reduced language barriers
- Faster customer response
- 24/7 automated assistance

Thus, the Bidirectional Encoder captures richer contextual information, while the decoder converts that representation into a meaningful target-language sequence.

Q4. BPTT and LSTM

A financial analytics company is developing a model to forecast stock prices using historical market data. Create an LSTM-based prediction system and explain how Backpropagation Through Time (BPTT) is used to train the network. Justify why LSTM is more suitable than a traditional RNN for this application.

Solution:

BPTT AND LSTM FOR STOCK PRICE FORECASTING


1. Industry Problem

A financial analytics company wants to forecast stock prices using historical market information. Stock prices are time-series data where the current value may depend on patterns from previous time steps.

The input can include:

- Open price
- High price
- Low price
- Closing price
- Trading volume
- Technical indicators

2. Proposed LSTM Architecture



3. Why LSTM?

Traditional RNNs can suffer from:

- Vanishing gradients
- Difficulty learning long-term dependencies
- Loss of information over long sequences

LSTM addresses these problems through a memory cell and gates.

4. LSTM Gates

An LSTM contains:

Forget Gate

Determines what old information should be discarded.

$$f_t = \sigma(W_f[h_{t-1}, x_t] + b_f)$$

Input Gate

Determines what new information should be stored.

$$i_t = \sigma(W_i[h_{t-1}, x_t] + b_i)$$

Candidate Memory

$$\tilde{C}_t = \tanh(W_c[h_{t-1}, x_t] + b_c)$$

Cell State

$$C_t = f_t C_{t-1} + i_t \tilde{C}_t$$

Output Gate

$$o_t = \sigma(W_o[h_{t-1}, x_t] + b_o)$$

The hidden state is:


$$h_t = o_t \tanh(C_t)$$

These mechanisms allow LSTM to preserve useful information over longer time periods.

5. BPTT

Backpropagation Through Time (BPTT) is used to train recurrent neural networks.

The LSTM is conceptually unfolded across time:




The prediction error is calculated using a loss function.

For regression, Mean Squared Error can be used:


$$MSE = (1/n) \sum (y_i - \hat{y}_i)^2$$

The error is propagated backward through the sequence.



An optimizer such as Adam updates the model parameters.

6. Training Process



The data must be split chronologically rather than randomly to avoid future-data leakage.

7. Evaluation

Useful regression metrics include:

- MAE
- MSE
- RMSE
- MAPE

The model can also be evaluated on directional accuracy:

Predicted direction vs Actual direction

8. Why LSTM is More Suitable Than Traditional RNN

Traditional RNN	LSTM
Can suffer from vanishing gradients	Better handles gradient problems
Weak long-term memory	Stronger long-term memory
Simple architecture	Memory cell + gates
Difficult for long sequences	Better for long sequences

9. Industry Considerations

Stock prices are influenced by many external factors such as:

- Market sentiment
- Economic events
- Company announcements
- Interest rates
- Global events

Therefore, an LSTM should not be considered a guaranteed stock-price prediction system. It should be validated using historical out-of-sample data and used as an analytical forecasting tool.

Q5. Integrated Deep Learning Solution

A media streaming company plans to automatically generate captions for uploaded videos to improve accessibility. Design an integrated deep learning solution that combines Transfer Learning for visual feature extraction with an LSTM-based Encoder–Decoder architecture for caption generation. Explain the role of each component and justify how the complete system produces accurate captions.

Solution:

INTEGRATED DEEP LEARNING SOLUTION FOR AUTOMATIC VIDEO CAPTIONING

1. Industry Problem

A media streaming company wants to automatically generate captions for uploaded videos.


The system must understand both:


1. **What appears in the video**
2. **What is happening over time**

A single image classifier is insufficient because videos contain temporal information.

Therefore, an integrated architecture combining Transfer Learning + LSTM Encoder–Decoder is proposed.

2. Overall Architecture






3. Component 1 – Transfer Learning

A pre-trained CNN such as ResNet50 or EfficientNet is used as the visual feature extractor. The CNN has already learned useful visual features from a large image dataset.

For every video frame:



The final classification layer is removed because we do not want to classify the frame. Instead, we want to extract its visual representation.

4. Why Transfer Learning?

Training a CNN from scratch would require a huge image/video dataset.

Transfer learning:

- Reduces training time
- Reduces data requirements
- Provides strong visual features
- Improves generalization
- Reduces computational requirements

Initially, the CNN backbone can be frozen. Later, selected deeper layers can be fine-tuned using a small learning rate.

5. Component 2 – LSTM Encoder

The extracted features from successive video frames form a sequence:

Frame 1 → Feature 1

Frame 2 → Feature 2

Frame 3 → Feature 3


...

Frame N → Feature N

The LSTM encoder processes this sequence.

It learns temporal relationships such as:

Person enters room




This temporal understanding is important for generating meaningful captions.

6. Component 3 – LSTM Decoder

The decoder converts the learned video representation into natural language.

Example:




The decoder predicts one word at a time.

7. Attention Mechanism

An attention mechanism can be added to improve caption quality.

Instead of forcing the decoder to rely on one fixed context vector, attention allows it to focus on different frame features while generating different words.

For example:




This can improve temporal and semantic alignment between video content and generated words.


8. Training

The model can be trained using video-caption datasets such as:

- MSR-VTT
- MSVD
- ActivityNet Captions

Training process:





Teacher forcing can be used during training, where the correct previous word is provided to the decoder.

9. Performance Evaluation


Caption quality can be evaluated using:

- BLEU
- ROUGE-L
- METEOR
- CIDEr
- Human evaluation

CIDEr is particularly useful for image/video captioning because it evaluates similarity between generated and reference captions.

10. Industry Deployment

A production pipeline could work as follows:



The generated captions can be stored as subtitle files such as SRT or WebVTT.

11. Complete System Advantages

The integrated system combines the strengths of both approaches:

Component	Main Role
Transfer Learning CNN	Understands visual content
Frame Extraction	Converts video into manageable image sequences
LSTM Encoder	Learns temporal relationships
Attention	Focuses on relevant frames/features
LSTM Decoder	Generates natural-language captions
BPTT	Trains temporal dependencies
Evaluation Metrics	Measures caption quality

12. Final Analysis

Transfer learning provides strong visual representations without requiring the CNN to learn everything from scratch. The LSTM encoder captures how visual information changes across time, while the decoder converts this information into a natural-language sequence. Therefore:

Visual Understanding

+

Temporal Understanding

+

Language Generation

=

Automatic Video Captioning

The complete architecture can provide scalable automated caption generation for media platforms, improving accessibility for users who are deaf or hard of hearing and making video content easier to search and understand.

Overall Industry-Level Tool Stack

The following technologies can be used to implement the proposed solutions:

Area	Tools / Technologies
Programming	Python
Deep Learning	PyTorch / TensorFlow
Computer Vision	OpenCV
Pre-trained Models	torchvision / TensorFlow Hub / Hugging Face
NLP	Hugging Face Transformers / NLTK
GPU Acceleration	CUDA
Experiment Tracking	TensorBoard / MLflow
Data Processing	NumPy / Pandas
Visualization	Matplotlib
Deployment	Flask / FastAPI / Docker
Hardware	NVIDIA GPU / Cloud GPU

RUBRICS FOR CALCULATION:

Criteria	Weightage	Level 4 – Exemplary	Level 3 – Proficient	Level 2 – Developing	Level 1 – Beginning
Understanding of Industry Problem	20%	Comprehensive understanding of real-world problem and constraints	Good understanding with minor gaps in requirements	Basic understanding ; some requirements unclear	Poor understanding ; misinterprets problem context
Application of Engineering Concepts	25%	Applies advanced and relevant concepts effectively to solve the problem	Applies appropriate concepts with minor errors	Limited application of concepts	Incorrect or no application of concepts
Solution Design	25%	Innovative, practical, and feasible solution aligned with	Logical and feasible solution with minor gaps	Basic solution with limited feasibility	Unclear or impractical solution

		industry standards			
Use of Tools	20%	Effective use of modern tools, technologies, and real data	Appropriate use of tools with correct implementation	Limited or inconsistent use of tools	Minimal or incorrect use of tools
Analysis	10%	Thorough analysis with validated results and performance evaluation	Good analysis with reasonable validation	Basic analysis with limited validation	Poor or no analysis